



PROJECT

Generate Faces

A part of the Deep Learning Nanodegree Foundation Program

PROJECT REVIEW

CODE REVIEW

NOTES

SHARE YOUR ACCOMPLISHMENT!  

Meets Specifications

This is an awesome submission. The faces generated are quite clear and realistic. I have left my views, suggestions and comments below to your reference.

In order to get good performance with image generation, there are several different implementations and tricks, which you can check out from <https://github.com/soumith/ganhacks>

Here are some other important resources for GAN:

<http://www.araya.org/archives/1183> for GAN stability.

<https://github.com/yihui-he/GAN-MNIST>, <https://github.com/carpdm20/DCGAN-tensorflow> for DCGAN.

<https://medium.com/@ageitgey/abusing-generative-adversarial-networks-to-make-8-bit-pixel-art-e45d9b96cee7>

It can be seen that a lot of hard work has been put into this.

Congratulations on successfully completing Project 5 and the deep learning nanodegree! :)

Required Files and Tests

The project submission contains the project notebook, called "dLnd_face_generation.ipynb".

All the unit tests in project have passed.

Build the Neural Network

The function `model_inputs` is implemented correctly.

The function `discriminator` is implemented correctly.

Good work using Batch Normalization.

Nice work using Leaky RELUs, which allow a small, non-zero gradient when the unit is not active.

Try using different values of alpha between 0.08 and 0.15 and compare your results.

The function `generator` is implemented correctly.

Good work using Batch Normalization.

Nice work using Leaky RELUs which allow a small, non-zero gradient when the unit is not active.

Since tanh is the last layer of the generator output, normalizing the input images is a good step.

You should use different values of alpha between 0.08 and 0.15 and compare your results.
Brilliant Work!!

The function `model_loss` is implemented correctly.

Good work using smoothing as it prevents discriminator from being from being too strong and to generalize in a better way better way

The function `model_opt` is implemented correctly.

Good Implementation!!

You have wrapped the training step with dependency to `tf.GraphKeys.UPDATE_OPS`, which is brilliant.

Neural Network Training

The function `train` is implemented correctly.

- It should build the model using `model_inputs`, `model_loss`, and `model_opt`.
- It should show output of the `generator` using the `show_generator_output` function

Good job keeping `batch_z` between -1 and 1.

Good work increasing the batch size by a factor of 2, i.e., `batch_images = batch_images * 2`, inside the inner for loop.

The parameters are set reasonable numbers.

Good work adjusting the hyper parameters.

Try using different values of learning rate between 0.0002 and 0.0008 and different values of Beta1 between 0.2 and 0.5 and compare your results.

The project generates realistic faces. It should be obvious that images generated look like faces.

The faces generated are quite clear and realistic. :)

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)

[Student FAQ](#)

[Reviewer Agreement](#)